

Imagine that one has a data matrix  $X \in \mathbb{R}^{n \times p}$  consisting of  $n$  observations, each with  $p$  features, as well as a response vector  $y \in \mathbb{R}^n$ . We want to build a model for  $y$  using the feature columns in  $X$ . In **ordinary least squares (OLS)**, one seeks a vector of coefficients  $\hat{\beta} \in \mathbb{R}^p$  such that

$$\hat{\beta} = \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \|y - X\beta\|_2^2.$$

In **non-negative least squares (NNLS)**, we seek a vector coefficients  $\hat{\beta} \in \mathbb{R}^p$  such that it minimizes  $\|y - X\beta\|_2^2$  subject to the additional requirement that each element of  $\hat{\beta}$  is non-negative.

There are a number of ways to perform NNLS in R. The first two methods come from Reference 1, while I came up with the third. (I'm not sharing the third way Reference 1 details because it claims that the method is buggy.)

Let's generate some fake data that we will use for the rest of the post:

```
set.seed(1)
n <- 100; p <- 10
x <- matrix(rnorm(n * p), nrow = n)
y <- x %*% matrix(rep(c(1, -1), length.out = p), ncol = 1) + rnorm(n)
```

### **Method 1: the nnls package**

```
library(nnls)
mod1 <- nnls(x, y)
mod1$x
# [1] 0.9073423 0.0000000 1.2971069 0.0000000 0.9708051
# [6] 0.0000000 1.2002310 0.0000000 0.3947028 0.0000000
```

### **Method 2: the glmnet package**

The `glmnet()` function solves the minimization problem

$$\hat{\beta} = \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \frac{1}{2n} \|y - X\beta\|_2^2 + \lambda \left[ \frac{1 - \alpha}{2} \|\beta\|_2^2 + \alpha \|\beta\|_1 \right],$$

where  $\alpha$  and  $\lambda$  are hyperparameters the user chooses. By setting  $\alpha = 1$  (the default) and  $\lambda = 0$ , `glmnet()` ends up solving the OLS problem. By setting `lower.limits = 0`, this forces the coefficients to be non-negative. We should also set `intercept = FALSE` so that we don't have an extraneous intercept term.

```
library(glmnet)
mod2 <- glmnet(x, y, lambda = 0, lower.limits = 0, intercept = FALSE)
coef(mod2)
# 11 x 1 sparse Matrix of class "dgCMatrix"
# s0
# (Intercept) .
# V1          0.9073427
# V2          .
# V3          1.2971070
# V4          .
# V5          0.9708049
# V6          .
# V7          1.2002310
# V8          .
# V9          0.3947028
# V10         .
```

### **Method 3: the bvl package**

NNLS is a special case of **bounded-variable least squares (BVLS)**, where instead of having constraints  $\beta_j \geq 0$  for each  $j = 1, \dots, p$ , one has constraints  $a_j \leq \beta_j \leq b_j$  for each  $j$ . BVLS is implemented in the `bvls()` function of the `bvls` package:

```
library(bvls)
mod3 <- bvls(x, y, bl = rep(0, p), bu = rep(Inf, p))
mod3$x
# [1] 0.9073423 0.0000000 1.2971069 0.0000000 0.9708051
# [6] 0.0000000 1.2002310 0.0000000 0.3947028 0.0000000
```

In the above, `bl` contains the lower limits for the coefficients while `bu` contains the upper limits for the coefficients.

References:

1. Things I Thought At One Point. [Three ways to do non-negative least squares in R.](#)